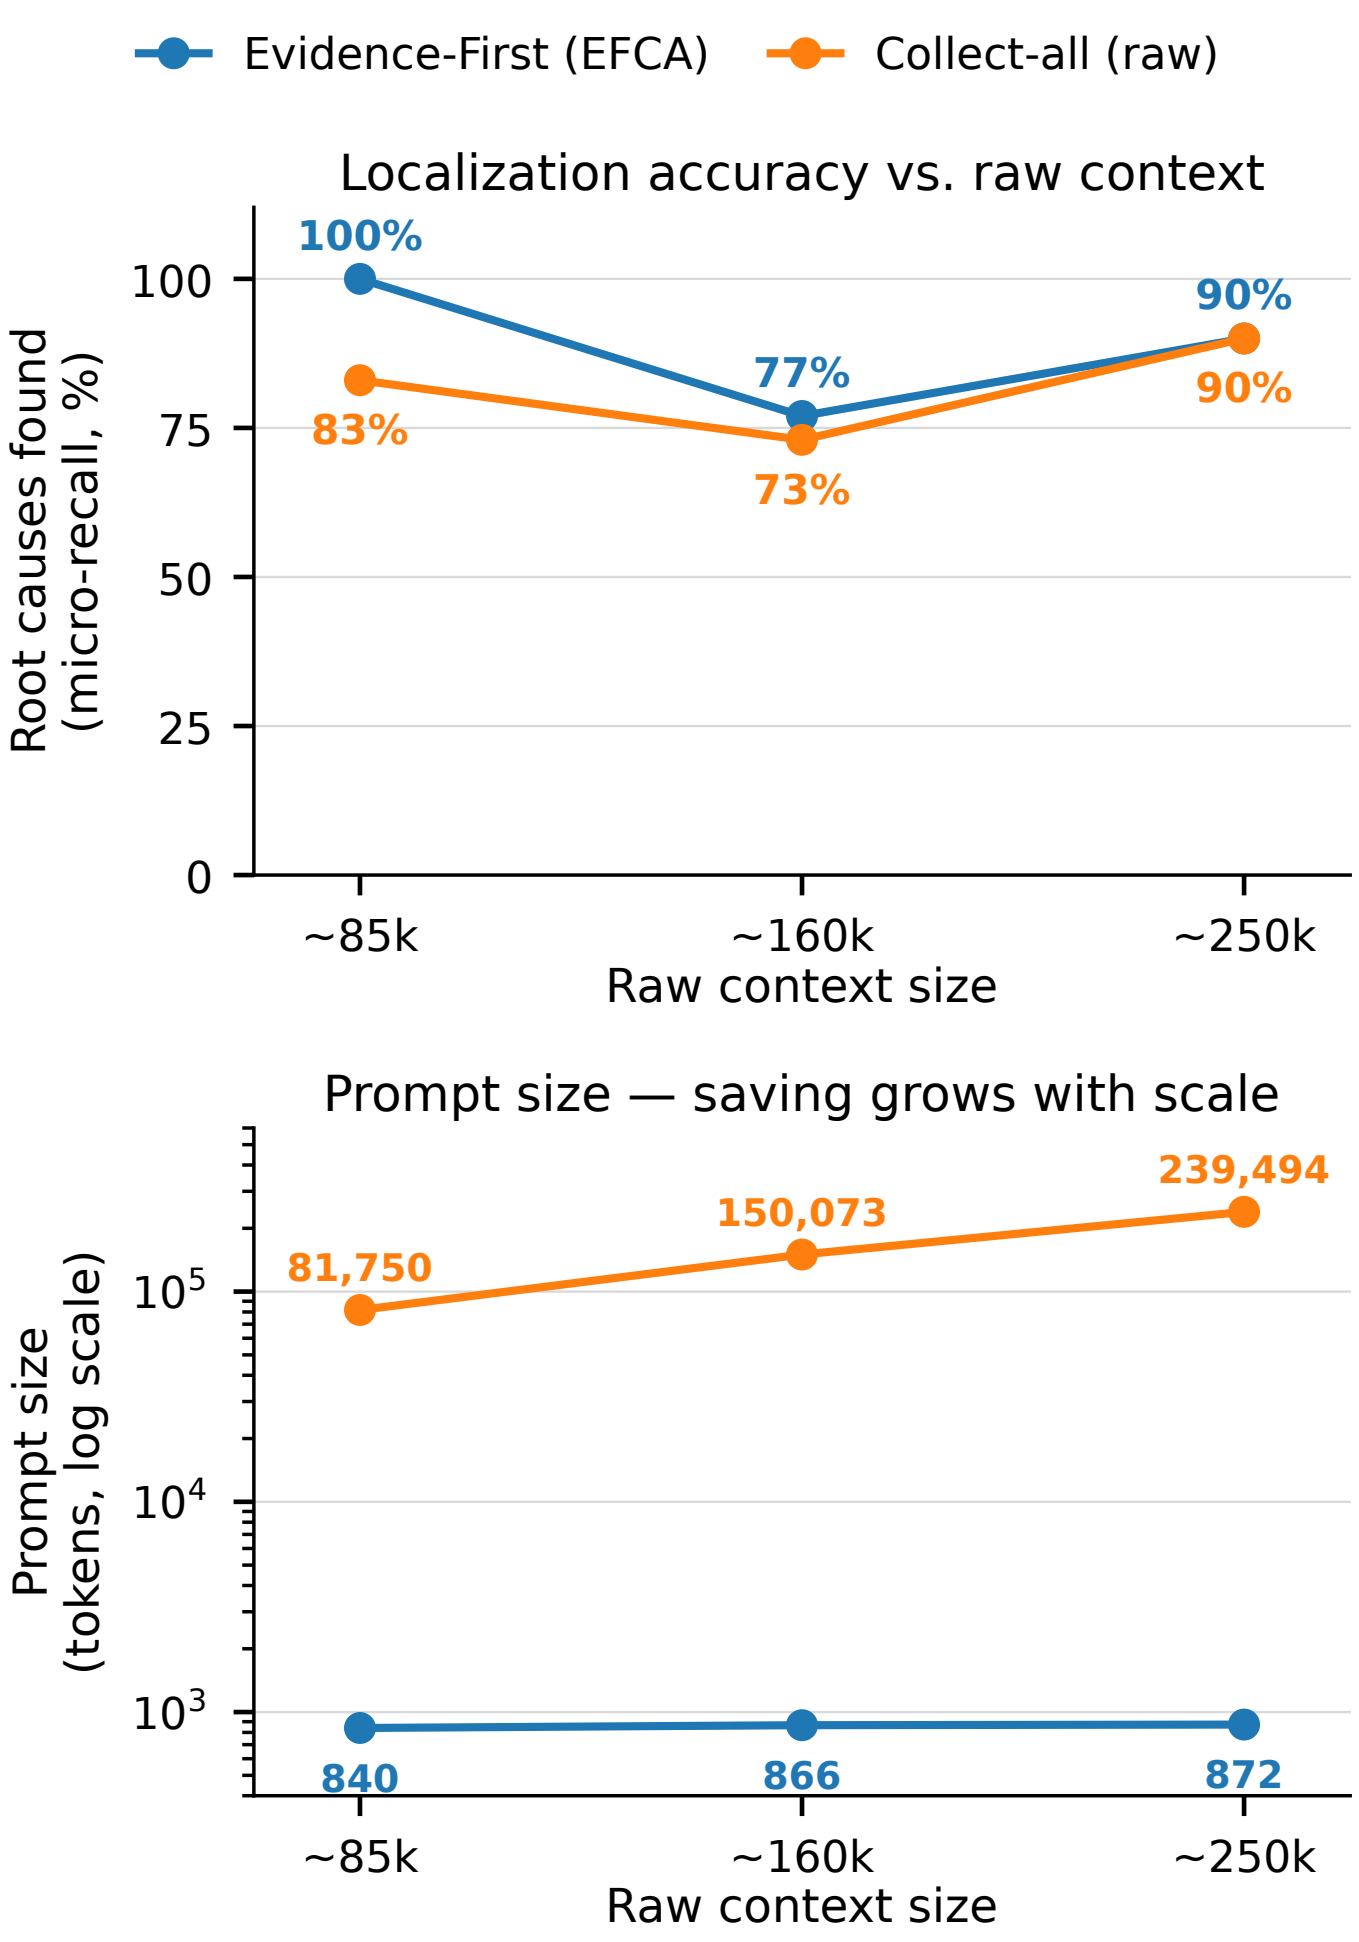
Evidence-First (EFCA)
Collect-all (raw)
Localization accuracy vs. raw context
Root causes found
(micro-recall, %)
100
75
50
25
0
100%
77%
90%
83%
73%
90%
~85k
~160k
~250k
Raw context size
Prompt size — saving grows with scale
Prompt size
(tokens, log scale)
10^5
10^4
10^3
239,494
150,073
81,750
840
866
872
~85k
~160k
~250k
Raw context size